

Contents

Confluent Platform GitOps with ArgoCD	2
Complete Guide to Declarative Kafka Management on Kubernetes	2
Table of Contents	2
Executive Summary	3
Key Technologies	3
Architecture Overview	3
Repository Structure	3
Prerequisites	4
Required Tools	4
AWS Permissions	4
Part 1: EKS Cluster Setup	4
1.1 Create EKS Cluster Configuration	4
1.2 Create the Cluster	5
1.3 Install EBS CSI Driver	5
1.4 Create Storage Class	6
1.5 Verify Cluster	6
Part 2: Install CFK Operator	6
2.1 Add Confluent Helm Repository	6
2.2 Create Operator Namespace	6
2.3 Install CFK Operator	6
2.4 Verify Installation	7
Part 3: Install and Configure ArgoCD	7
3.1 Create ArgoCD Namespace	7
3.2 Install ArgoCD	7
3.3 Wait for ArgoCD to be Ready	7
3.4 Expose ArgoCD UI (NodePort)	7
3.5 Get ArgoCD Access Details	7
3.6 Access ArgoCD	8
Part 4: Deploy Confluent Platform via GitOps	8
4.1 Create ArgoCD Application Manifest	8
4.2 Deploy the Application	8
4.3 Monitor Deployment	9
4.4 Verify Deployment	9
Part 5: Expose Services for External Access	9
5.1 Expose Control Center via NodePort	9
5.2 Get Node IP and Update Config	9
5.3 Commit and Push Changes	9
5.4 ArgoCD Auto-Syncs	10
5.5 Access Control Center	10
Part 6: Connector GitOps Demo	10
6.1 Create Connector Demo Repository	10
6.2 Add Kubernetes Deployment Manifests	10
6.3 Commit and Push	12
6.4 Deploy to ArgoCD	12
6.5 Register Repository for Branch Discovery	12
6.6 Create Topic	12

6.7 Verify Connector	12
Part 7: Branch-Based Configuration Management	12
7.1 Create Feature Branches	12
7.2 Branch Configuration Summary	14
7.3 Switch Branches via ArgoCD UI	14
7.4 Switch Branches via Command Line	14
7.5 Verify Branch Switch	14
Operational Commands Reference	15
ArgoCD Commands	15
Confluent Platform Commands	15
Topic Management	15
Troubleshooting Guide	16
ArgoCD Not Syncing	16
Connector in ERROR State	16
Pods Not Starting	16
Topic Not Created	16
Cleanup	17
Delete Connector Demo	17
Delete Confluent Platform	17
Delete ArgoCD	17
Delete EKS Cluster	17
Summary	17
Key Benefits	17
Resources	17

Confluent Platform GitOps with ArgoCD

Complete Guide to Declarative Kafka Management on Kubernetes

Date: January 28, 2026

Author: Confluent Federal

Version: 1.0

Table of Contents

1. Executive Summary
2. Architecture Overview
3. Prerequisites
4. Part 1: EKS Cluster Setup
5. Part 2: Install CFK Operator
6. Part 3: Install and Configure ArgoCD
7. Part 4: Deploy Confluent Platform via GitOps
8. Part 5: Expose Services for External Access
9. Part 6: Connector GitOps Demo
10. Part 7: Branch-Based Configuration Management
11. Operational Commands Reference
12. Troubleshooting Guide

13. Cleanup

Executive Summary

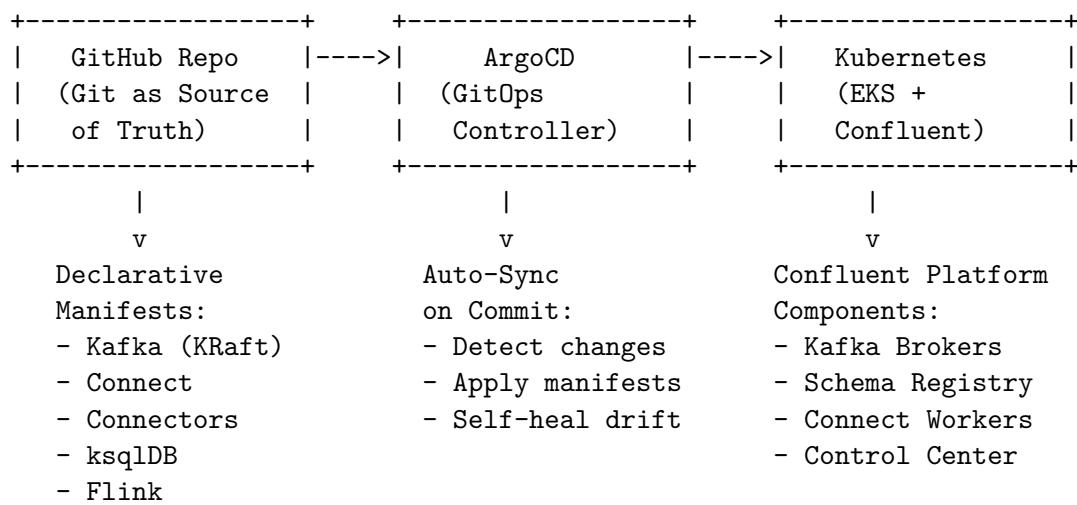
This white paper documents the complete implementation of a GitOps-based Confluent Platform deployment on AWS EKS using ArgoCD. The solution demonstrates:

- **Infrastructure as Code** - EKS cluster defined in YAML
- **GitOps Deployment** - All Confluent components managed via Git
- **Declarative Connectors** - Kafka Connect connectors defined as Kubernetes CRDs
- **Branch-Based Configuration** - Different environments/configs via Git branches
- **Self-Healing Infrastructure** - ArgoCD automatically reverts drift

Key Technologies

Technology	Version	Purpose
AWS EKS	1.31	Kubernetes cluster
Confluent Platform	8.1	Event streaming platform
CFK Operator	3.1	Kubernetes operator for CP
ArgoCD	Latest stable	GitOps continuous delivery
Kustomize	Built-in	Configuration management

Architecture Overview



Repository Structure

```
cfk-argocd-demo/
|-- argocd-application.yaml      # ArgoCD Application manifest
|-- eks-cluster-config.yaml     # EKS cluster definition
|-- base/
```

```
|   |-- kustomization.yaml      # Base kustomization
|   |-- confluent-platform.yaml # Core CP components
|   |-- connectors.yaml        # Declarative connectors
|   |-- ksqlDB-queries.yaml    # ksqlDB queries as code
|   |-- kstreams-app.yaml      # Kafka Streams app
|   +-- flink.yaml             # Flink resources (optional)
|-- overlays/
|   |-- dev/                   # Development overrides
|   +-- prod/                  # Production overrides
+-- docs/
    +-- CONFLUENT-PLATFORM-GITOPS-WHITEPAPER.md
```

Prerequisites

Required Tools

```
# AWS CLI
aws --version

# eksctl
eksctl version

# kubectl
kubectl version --client

# Helm
helm version

# GitHub CLI
gh --version
```

AWS Permissions

- EKS cluster creation
 - EC2 instance management
 - EBS volume provisioning
 - IAM role management
-

Part 1: EKS Cluster Setup

1.1 Create EKS Cluster Configuration

Create `eks-cluster-config.yaml`:

```
apiVersion: eksctl.io/v1alpha5
kind: ClusterConfig
```

```

metadata:
  name: cfk-argocd-demo
  region: us-east-1
  version: "1.31"
  tags:
    owner_email: your-email@company.com
    environment: demo

vpc:
  nat:
    gateway: Disable

managedNodeGroups:
  - name: cfk-argocd-nodes
    instanceType: m5.xlarge
    desiredCapacity: 3
    minSize: 3
    maxSize: 6
    privateNetworking: false
    volumeSize: 100
    labels:
      role: confluent

```

1.2 Create the Cluster

```

# Create EKS cluster (takes 15-20 minutes)
eksctl create cluster -f eks-cluster-config.yaml

```

1.3 Install EBS CSI Driver

```

# Install EBS CSI addon
eksctl create addon \
  --name aws-ebs-csi-driver \
  --cluster cfk-argocd-demo \
  --region us-east-1

# Get node role name
ROLE_NAME=$(aws iam list-roles \
  --query "Roles[?contains(RoleName, 'eksctl-cfk-argocd-demo-nodegroup')].RoleName" \
  --output text)

# Attach EC2 permissions
aws iam attach-role-policy \
  --role-name "$ROLE_NAME" \
  --policy-arn arn:aws:iam::aws:policy/AmazonEC2FullAccess

```

1.4 Create Storage Class

```
kubectl apply -f - <<EOF
kind: StorageClass
apiVersion: storage.k8s.io/v1
metadata:
  name: gp3
  annotations:
    storageclass.kubernetes.io/is-default-class: "true"
provisioner: ebs.csi.aws.com
parameters:
  type: gp3
  encrypted: "true"
volumeBindingMode: WaitForFirstConsumer
EOF
```

1.5 Verify Cluster

```
# Check nodes
kubectl get nodes

# Expected output:
# NAME                  STATUS   ROLES      AGE     VERSION
# ip-xxx-xxx-xxx-xxx...  Ready    <none>    5m      v1.31.x
# ip-xxx-xxx-xxx-xxx...  Ready    <none>    5m      v1.31.x
# ip-xxx-xxx-xxx-xxx...  Ready    <none>    5m      v1.31.x
```

Part 2: Install CFK Operator

2.1 Add Confluent Helm Repository

```
helm repo add confluentinc https://packages.confluent.io/helm
helm repo update
```

2.2 Create Operator Namespace

```
kubectl create namespace confluent-operator
```

2.3 Install CFK Operator

```
helm upgrade --install confluent-operator \
  confluentinc/confluent-for-kubernetes \
  --namespace confluent-operator \
  --set namespaced=false \
  --wait
```

2.4 Verify Installation

```
# Check operator pod
kubectl get pods -n confluent-operator
```

```
# Check CRDs installed
kubectl get crds | grep confluent
```

Expected CRDs: - kafkas.platform.confluent.io - kafkarestclasses.platform.confluent.io
- connects.platform.confluent.io - connectors.platform.confluent.io - schemaregistries.platform.confluent.io
- ksqldbplatform.confluent.io - controlcenters.platform.confluent.io

Part 3: Install and Configure ArgoCD

3.1 Create ArgoCD Namespace

```
kubectl create namespace argocd
```

3.2 Install ArgoCD

```
kubectl apply -n argocd \
-f https://raw.githubusercontent.com/argoproj/argo-cd/stable/manifests/install.yaml
```

3.3 Wait for ArgoCD to be Ready

```
kubectl wait --for=condition=ready pod \
-l app.kubernetes.io/name=argocd-server \
-n argocd \
--timeout=300s
```

3.4 Expose ArgoCD UI (NodePort)

```
kubectl patch svc argocd-server -n argocd \
-p '{"spec": {"type": "NodePort"}}'
```

3.5 Get ArgoCD Access Details

```
# Get NodePort
ARGOCD_PORT=$(kubectl get svc argocd-server -n argocd \
-o jsonpath='{.spec.ports[0].nodePort}')

# Get Node External IP
NODE_IP=$(kubectl get nodes \
-o jsonpath='{.items[0].status.addresses[?(@.type=="ExternalIP")].address}')

# Get initial admin password
ARGOCD_PASSWORD=$(kubectl -n argocd get secret argocd-initial-admin-secret \
-o jsonpath="{.data.password}" | base64 -d)
```

```
echo "ArgoCD URL: https://$NODE_IP:$ARGOCD_PORT"
echo "Username: admin"
echo "Password: $ARGOCD_PASSWORD"
```

3.6 Access ArgoCD

Open the URL in your browser and login with the credentials above.

Part 4: Deploy Confluent Platform via GitOps

4.1 Create ArgoCD Application Manifest

Create argocd-application.yaml:

```
apiVersion: argoproj.io/v1alpha1
kind: Application
metadata:
  name: confluent-platform
  namespace: argocd
spec:
  project: default
  source:
    repoURL: https://github.com/confluentfederal/cfk-argocd-demo.git
    targetRevision: HEAD
    path: overlays/prod
  destination:
    server: https://kubernetes.default.svc
    namespace: confluent
  syncPolicy:
    automated:
      prune: true
      selfHeal: true
  syncOptions:
    - CreateNamespace=true
    - PrunePropagationPolicy=foreground
  retry:
    limit: 5
    backoff:
      duration: 5s
      factor: 2
      maxDuration: 3m
```

4.2 Deploy the Application

```
kubectl apply -f argocd-application.yaml
```

4.3 Monitor Deployment

```
# Watch ArgoCD application status
kubectl get application confluent-platform -n argocd -w

# Watch pods coming up
kubectl get pods -n confluent -w
```

4.4 Verify Deployment

```
# Check all Confluent components
kubectl get kafka,connect,schemaregistry,ksqldb,controlcenter -n confluent
```

Expected Components: - Kafka (3 brokers in KRaft mode) - Schema Registry - Kafka Connect - ksqlDB - Control Center

Part 5: Expose Services for External Access

5.1 Expose Control Center via NodePort

Update base/confluent-platform.yaml to add external access:

```
apiVersion: platform.confluent.io/v1beta1
kind: ControlCenter
metadata:
  name: controlcenter
  namespace: confluent
spec:
  # ... existing config ...
  externalAccess:
    type: nodePort
    nodePort:
      host: <NODE_EXTERNAL_IP>
      nodePortOffset: 30021
```

5.2 Get Node IP and Update Config

```
# Get external IP
NODE_IP=$(kubectl get nodes \
-o jsonpath='{.items[0].status.addresses[?(@.type=="ExternalIP")].address}')

echo "Use this IP in externalAccess.nodePort.host: $NODE_IP"
```

5.3 Commit and Push Changes

```
cd /path/to/cfk-argocd-demo
git add base/confluent-platform.yaml
git commit -m "Add external access for Control Center"
git push origin main
```

5.4 ArgoCD Auto-Syncs

ArgoCD detects the Git change and automatically applies it. Monitor in UI or:

```
kubectl get application confluent-platform -n argocd \
-o jsonpath='{.status.sync.status}'
```

5.5 Access Control Center

```
# Get the NodePort
CC_PORT=$(kubectl get svc controlcenter-bootstrap -n confluent \
-o jsonpath='{.spec.ports[0].nodePort}')

echo "Control Center: http://$NODE_IP:$CC_PORT"
```

Part 6: Connector GitOps Demo

This section demonstrates managing Kafka Connect connectors via Git with branch-based configuration.

6.1 Create Connector Demo Repository

```
# Clone baseline connector code
git clone https://github.com/confluentfederal/kafka-connect-netty-source.git \
kafka-connect-netty-source-argo-example

cd kafka-connect-netty-source-argo-example

# Initialize fresh git history
rm -rf .git
git init
git add .
git commit -m "Initial commit - Netty Source Connector baseline"

# Create GitHub repo and push
gh repo create confluentfederal/kafka-connect-netty-source-argo-example \
--public --source=. --remote=origin --push
```

6.2 Add Kubernetes Deployment Manifests

Create `deploy/base/connector.yaml`:

```
apiVersion: platform.confluent.io/v1beta1
kind: Connector
metadata:
  name: datagen-gitops-demo
  namespace: confluent
spec:
  class: io.confluent.kafka.connect.datagen.DatagenConnector
```

```

taskMax: 1
connectClusterRef:
  name: connect
configs:
  kafka.topic: "gitops-pageviews"
  quickstart: "pageviews"
  key.converter: "org.apache.kafka.connect.storage.StringConverter"
  value.converter: "io.confluent.connect.avro.AvroConverter"
  value.converter.schema.registry.url: "http://schemaregistry.confluent.svc.cluster.local:8081"
  max.interval: "1000"
  iterations: "10000000"

```

Create deploy/base/kustomization.yaml:

```

apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

resources:
  - connector.yaml

commonLabels:
  app.kubernetes.io/name: datagen-gitops-demo
  app.kubernetes.io/managed-by: argocd

```

Create deploy/argocd-application.yaml:

```

apiVersion: argoproj.io/v1alpha1
kind: Application
metadata:
  name: netty-connector
  namespace: argocd
  finalizers:
    - resources-finalizer.argocd.argoproj.io
spec:
  project: default
  source:
    repoURL: https://github.com/confluentfederal/kafka-connect-netty-source-argo-example.git
    targetRevision: main
    path: deploy/base
  destination:
    server: https://kubernetes.default.svc
    namespace: confluent
  syncPolicy:
    automated:
      prune: true
      selfHeal: true

```

6.3 Commit and Push

```
git add deploy/
git commit -m "Add Kubernetes deploy manifests for ArgoCD"
git push origin main
```

6.4 Deploy to ArgoCD

```
kubectl apply -f deploy/argocd-application.yaml
```

6.5 Register Repository for Branch Discovery

To enable branch dropdown in ArgoCD UI:

```
kubectl apply -f - <<EOF
apiVersion: v1
kind: Secret
metadata:
  name: repo-netty-connector
  namespace: argocd
  labels:
    argocd.argoproj.io/secret-type: repository
stringData:
  type: git
  url: https://github.com/confluentfederal/kafka-connect-netty-source-argo-example.git
EOF
```

6.6 Create Topic

```
kubectl exec -n confluent kafka-0 -- kafka-topics \
  --bootstrap-server localhost:9092 \
  --create --topic gitops-pageviews \
  --partitions 3 --replication-factor 1
```

6.7 Verify Connector

```
kubectl get connector datagen-gitops-demo -n confluent
```

Part 7: Branch-Based Configuration Management

7.1 Create Feature Branches

Users Data Branch

```
git checkout -b feature/users-data
```

```
# Update deploy/base/connector.yaml
cat > deploy/base/connector.yaml <<EOF
apiVersion: platform.confluent.io/v1beta1
kind: Connector
```

```

metadata:
  name: datagen-gitops-demo
  namespace: confluent
spec:
  class: io.confluent.kafka.connect.datagen.DatagenConnector
  taskMax: 2
  connectClusterRef:
    name: connect
configs:
  kafka.topic: "gitops-users"
  quickstart: "users"
  key.converter: "org.apache.kafka.connect.storage.StringConverter"
  value.converter: "io.confluent.connect.avro.AvroConverter"
  value.converter.schema.registry.url: "http://schemaregistry.confluent.svc.cluster.local:8081"
  max.interval: "2000"
  iterations: "10000000"
EOF

git add .
git commit -m "Generate users data instead of pageviews"
git push -u origin feature/users-data

```

High-Throughput Orders Branch

```

git checkout main
git checkout -b feature/orders-high-throughput

# Update deploy/base/connector.yaml
cat > deploy/base/connector.yaml <<EOF
apiVersion: platform.confluent.io/v1beta1
kind: Connector
metadata:
  name: datagen-gitops-demo
  namespace: confluent
spec:
  class: io.confluent.kafka.connect.datagen.DatagenConnector
  taskMax: 4
  connectClusterRef:
    name: connect
configs:
  kafka.topic: "gitops-orders"
  quickstart: "orders"
  key.converter: "org.apache.kafka.connect.storage.StringConverter"
  value.converter: "io.confluent.connect.avro.AvroConverter"
  value.converter.schema.registry.url: "http://schemaregistry.confluent.svc.cluster.local:8081"
  max.interval: "100"
  iterations: "10000000"
EOF

```

```

git add .
git commit -m "Generate high-volume orders data"
git push -u origin feature/orders-high-throughput
git checkout main

```

7.2 Branch Configuration Summary

Branch	Topic	Data Type	Tasks	Rate
main	gitops-pageviews	Pageviews	1	1s
feature/users-data	gitops-users	Users	2	2s
feature/orders-high-throughput	high-throughput	Orders	4	100ms

7.3 Switch Branches via ArgoCD UI

1. Open ArgoCD UI
2. Click on netty-connector application
3. Click **APP DETAILS** then **EDIT**
4. Change **TARGET REVISION** dropdown to desired branch
5. Click **SAVE**
6. ArgoCD auto-syncs the new configuration

7.4 Switch Branches via Command Line

```

# Switch to users data
kubectl patch application netty-connector -n argocd \
--type merge -p '{"spec":{"source":{"targetRevision":"feature/users-data"}}}'

# Switch to high-throughput orders
kubectl patch application netty-connector -n argocd \
--type merge -p '{"spec":{"source":{"targetRevision":"feature/orders-high-throughput"}}}'

# Switch back to main
kubectl patch application netty-connector -n argocd \
--type merge -p '{"spec":{"source":{"targetRevision":"main"}}}'

```

7.5 Verify Branch Switch

```

# Check current branch
kubectl get application netty-connector -n argocd \
-o jsonpath='{.spec.source.targetRevision}'

# Check connector configuration
kubectl get connector datagen-gitops-demo -n confluent \
-o jsonpath='Topic: {.spec.configs.kafka\\.topic}, Tasks: {.spec.taskMax}'

```

Operational Commands Reference

ArgoCD Commands

Command	Description
kubectl get application -n argocd	List all applications
kubectl get application <name> -n argocd -o yaml	View application details
kubectl patch application <name> -n argocd --type merge -p '{...}'	Update application
kubectl annotate application <name> -n argocd argocd.argoproj.io/refresh=hard	Force refresh
kubectl patch application <name> -n argocd --type merge -p '{"operation": null}'	Clear stuck operation

Confluent Platform Commands

Command	Description
kubectl get kafka,connect,schemaregistry -n confluent	List CP components
kubectl get connector -n confluent	List connectors
kubectl describe connector <name> -n confluent	Connector details
kubectl logs -l app=connect -n confluent	Connect worker logs

Topic Management

Command	Description
kubectl exec kafka-0 -n confluent -- kafka-topics --list --bootstrap-server localhost:9092	List topics
kubectl exec kafka-0 -n confluent -- kafka-topics --create --topic <name> --bootstrap-server localhost:9092	Create topic

Command	Description
kubectl exec kafka-0 -n confluent -- kafka-console-consumer --topic <name> --bootstrap-server localhost:9092 --from-beginning	Consume messages

Troubleshooting Guide

ArgoCD Not Syncing

```
# Force hard refresh
kubectl annotate application confluent-platform -n argocd \
    argocd.argoproj.io/refresh=hard --overwrite

# Clear stuck operation
kubectl patch application confluent-platform -n argocd \
    --type merge -p '{"operation": null}'

# Check sync status
kubectl get application confluent-platform -n argocd \
    -o jsonpath='{.status.sync.status} {.status.health.status}'
```

Connector in ERROR State

```
# Get error message
kubectl get connector <name> -n confluent \
    -o jsonpath='{.status.conditions[*].message}'

# Check Connect logs
kubectl logs -l app=connect -n confluent --tail=100
```

Pods Not Starting

```
# Describe pod for events
kubectl describe pod <pod-name> -n confluent

# Check PVC binding
kubectl get pvc -n confluent

# Check resource limits
kubectl describe nodes | grep -A 5 "Allocated resources"
```

Topic Not Created

```
# Manually create topic
kubectl exec kafka-0 -n confluent -- kafka-topics \
```

```
--bootstrap-server localhost:9092 \
--create --topic <topic-name> \
--partitions 3 --replication-factor 1
```

Cleanup

Delete Connector Demo

```
kubectl delete application netty-connector -n argocd
kubectl delete secret repo-netty-connector -n argocd
```

Delete Confluent Platform

```
kubectl delete application confluent-platform -n argocd
```

Delete ArgoCD

```
kubectl delete namespace argocd
```

Delete EKS Cluster

```
eksctl delete cluster --name cfk-argocd-demo --region us-east-1
```

Summary

This implementation demonstrates a complete GitOps workflow for Confluent Platform:

1. **Infrastructure** - EKS cluster created via eksctl
2. **Operators** - CFK and ArgoCD installed via Helm/kubectl
3. **Platform** - Confluent Platform deployed via ArgoCD from Git
4. **Connectors** - Declarative connector management with branch-based configs
5. **Operations** - Self-healing, audit trail, rollback via Git

Key Benefits

- **Single Source of Truth** - All configuration in Git
 - **Audit Trail** - Git history shows all changes
 - **Rollback** - Revert to any previous state via Git
 - **Self-Healing** - ArgoCD reverts manual changes
 - **Multi-Environment** - Branches for dev/staging/prod
-

Resources

- **Main Repository:** <https://github.com/confluentfederal/cfk-argocd-demo>
- **Connector Demo:** <https://github.com/confluentfederal/kafka-connect-netty-source-argo-example>
- **CFK Documentation:** <https://docs.confluent.io/operator/current/overview.html>

- **ArgoCD Documentation:** <https://argo-cd.readthedocs.io/>
-

Document generated: January 28, 2026